LawAI

Building India's Harvey AI

Full Technical, Business & Market Report

A rigorous 9-phase evaluation of an AI legal copilot
for India's 1.7M+ advocates and 55M+ pending cases

Prepared by	Jatin Sikka
Date	April 2026
Status	Pre-Seed / Planning
Contact	jatinsikka30@gmail.com
GitHub	github.com/jatinsikka

Confidential — For internal use and potential co-founder discussions only.

Contents

1 Phase 1: Market Reality Check **3**
- 1.1 Is This Viable in India Today? . . . 3
  - 1.1.1 The Numbers That Matter . . . 3
  - 1.1.2 The Honest Assessment of Digitization . . . 3
  - 1.1.3 Law Fragmentation: Why This is Harder Than the US . . . 3
- 1.2 Who Are the Real Customers? . . . 4
- 1.3 Top 3 Use Cases Worth Building First . . . 4
- 1.4 Realistic Market Sizing . . . 5
- 1.5 Why This Will / Will Not Work . . . 5

2 Phase 2: Product Definition **6**
- 2.1 What the MVP Does (and Doesn't Do) . . . 6
- 2.2 What Makes This 10x Better . . . 6
- 2.3 The Magic Moment . . . 6
- 2.4 User Journey . . . 7
- 2.5 Real-World Usage Scenarios . . . 7

3 Phase 3: Data Strategy **8**
- 3.1 Data Sources and Accessibility . . . 8
- 3.2 Data Pipeline Design . . . 8
- 3.3 RAG Architecture (High-Level) . . . 9
- 3.4 Data Risks . . . 9

4 Phase 4: AI/ML System Design **10**
- 4.1 Base Model Strategy . . . 10
- 4.2 Technology Stack . . . 11
- 4.3 Key AI Capabilities . . . 11
- 4.4 Evaluation: How We Measure Correctness . . . 11
  - 4.4.1 IndiaLaw Bench (Proprietary Benchmark) . . . 11
  - 4.4.2 Hallucination Reduction Strategy . . . 12
- 4.5 Biggest Technical Risks . . . 12

5 Phase 5: Product UX **13**
- 5.1 Interface Decision: Chat + Document Viewer . . . 13
- 5.2 Key Screens . . . 13
- 5.3 Trust-Building Features (Critical for Lawyers) . . . 13

6 Phase 6: Go-to-Market Strategy **14**
- 6.1 Who We Sell to First . . . 14
- 6.2 Getting the First 10 Customers . . . 14
- 6.3 Pricing Model . . . 14

7 Phase 7: Risks & Failure Modes **16**

8 Phase 8: Execution Plan **17**
- 8.1 Week 1–2: Proof of Concept . . . 17
- 8.2 Month 1: Working Prototype . . . 17
- 8.3 Month 3: Beta-Ready Product . . . 17
- 8.4 Month 6: Paid Launch . . . 17
- 8.5 What NOT to Build Early . . . 17

9 Phase 9: Final Verdict **18**
- 9.1 Should We Build This? . . . 18
- 9.2 What Version of This Idea Actually Wins? . . . 18
- 9.3 The ONE Wedge That Dominates . . . 18
- 9.4 The Decisive Metric . . . 19

10 Appendix A: Competitive Intelligence **20**

10.1 Harvey AI ($11B Valuation) 20
10.2 Litera (Lito AI Agent) 20
10.3 Indian Competitors 20
11 Appendix B: Cost Budget **21**
11.1 MVP Build (Months 1–6) 21
11.2 Post-Launch Unit Economics 21

Phase 1: Market Reality Check

Is This Viable in India Today?

Yes, but with brutal caveats. The opportunity is massive and the timing is right, but execution must account for India-specific realities that most AI startups ignore.

The Numbers That Matter

Metric	Value	Source
Registered advocates	1.7M+ (growing 50K/yr)	Bar Council of India
Pending court cases	55.8 million (March 2026)	NJDG
Cases pending 5+ years	17.2 million	NJDG
India Legal AI market	$29.5M → $106M by 2030	Grand View Research
India Legal Tech market	$464M → $1.25B by 2030	Grand View Research
Legal tech funding (2025)	+781% YoY	Tracxn
GDP lost to judicial delays	1.5–2% annually	CSM / Economic Survey
Court records digitized	Only 5.9% of legacy records	PIB / e-Courts
Central Acts in force	∼891	India Code
District courts with e-pay	Only 33.3%	e-Courts project

The Honest Assessment of Digitization

Reality Check: India's Legal System is Mostly Offline

Only 5.9% of legacy court records have been digitized. Only 33.3% of district courts have e-payment for court fees. The Supreme Court has flagged "alarming" trends of lawyers submitting AI-generated fake citations (Feb 2026). Most lawyers in tier-2/3 cities still rely on physical law books, word-of-mouth, and Indian Kanoon's basic search. **This is both the obstacle and the opportunity.**

Law Fragmentation: Why This is Harder Than the US

- **891 Central Acts** + thousands of State-level statutes (each state has its own)
- **Concurrent List:** 52 subjects where both Centre and State can legislate — conflicting laws are common
- **Personal laws:** Hindu, Muslim, Christian, Parsi — each with different rules on marriage, inheritance, succession
- **25 High Courts** with different procedural rules and local precedent
- **Subordinate legislation:** Thousands of rules, regulations, notifications, circulars
- **Language:** Judgments in English (SC, most HCs) but also Hindi, regional languages in lower courts

Harvey only had to handle US federal + 50 state laws in one language. We handle a fundamentally more complex system.

Who Are the Real Customers?

Segment	Size	Pain Point	Willingness to Pay	Priority
Solo practitioners (tier-2/3)	~1.2M	Spend hours on research, no tools, use Indian Kanoon or nothing	Low (INR 500–999/mo)	**HIGH**
Mid-size litigation firms (10–50 lawyers)	~5,000 firms	Junior associates waste 40–60% of time on research/drafting	Medium (INR 2K–5K/seat/mo)	**HIGH**
In-house legal (startups/SMEs)	~50,000 teams	Compliance tracking, contract review, regulatory updates	Medium-High	Medium
Large law firms (Tier-1)	~100 firms	Already evaluating Harvey, Jhana; need India-specific depth	High ($50–100/seat/mo)	Low (later)
Law students	1.5M enrolled	Research for moots, internships	Very Low	Low

Ideal Customer Profile (ICP)

Primary: Mid-size litigation firms (10–50 lawyers) in Delhi, Mumbai, Bangalore, Hyderabad, Chennai. They bill enough to pay INR 4,999/mo for a team plan, have junior associates who waste the most time on research, and the senior partner can see the ROI immediately.

Secondary: Solo practitioners in tier-2/3 cities. Massive volume (1.2M+), low individual spend, but a free tier converts them into a distribution moat. They tell other lawyers.

Top 3 Use Cases Worth Building First

1. **Legal Research with Verified Citations.** "Find all SC/HC cases on anticipatory bail under Section 438 CrPC where the accused was a public servant." Today this takes a junior associate 4–8 hours on Manupatra/SCC. We do it in 2 minutes with verified, ranked citations.

2. **Case Summarization + Citation Status.** Upload a 50-page judgment → get a structured summary (facts, issues, holdings, ratio) + whether it's still good law (not overruled/distinguished). Lawyers do this manually before every hearing.

3. **Statute–Case Interplay.** "How have courts interpreted Section 138 of the Negotiable Instruments Act?" This requires cross-referencing statutes with case law — exactly what

the citation graph enables.

Realistic Market Sizing

Honest TAM/SAM/SOM — Not Inflated

- **TAM:** 1.7M advocates × INR 12,000/yr = INR 2,040Cr (∼$245M/yr)
- **SAM:** 200K advocates who are digitally active and can pay = INR 240Cr (∼$29M/yr)
- **SOM (Year 1):** 2,000 paying users × INR 999/mo avg = INR 2.4Cr/yr (∼$290K/yr)
- **SOM (Year 3):** 15,000 paying users × INR 1,500/mo avg = INR 27Cr/yr (∼$3.2M/yr)

This is real money, not fantasy. Jhana.ai is at INR 2.58Cr revenue after 3+ years. We need to beat that trajectory.

Why This Will / Will Not Work

Why It Will Work	Why It Might Not
55.8M pending cases = desperate need for efficiency	Only 5.9% digitization; most lawyers are offline
Incumbents (Manupatra, SCC) have dated UX and bolt-on AI	Indian lawyers are price-sensitive; INR 999/mo is a real decision
No Shepard's Citations equivalent exists for India	Fragmentation (891+ acts, 25 HCs, personal laws) makes data quality hard
Court judgments are public domain (Eastern Book Co v DB Modak, 2008)	Trust gap: SC flagged fake AI citations in Feb 2026
Legal tech funding surged 781% in 2025	Jhana.ai (Harvard, $1.8M funded) has 3-year head start
Lawyers lose 40–60% of time to research/drafting	44% of attorneys resist new technology adoption

Phase 2: Product Definition

What the MVP Does (and Doesn't Do)

MVP Does	MVP Does NOT
AI-powered legal research chat	Contract drafting / generation
Verified citations to real SC/HC judgments	Multi-language support (Hindi comes later)
Case summarization (upload judgment → structured summary)	e-Filing or court integration
Citation status (followed/distinguished/overruled)	Practice management (billing, CRM)
Statute–case cross-reference	Predictive case outcome analysis
Free tier (10 queries/day)	Offline mode

What Makes This 10x Better

Capability	ChatGPT	Manupatra / SCC Online	LawAI
Legal research	Hallucinates Indian cases confidently	Boolean search; you find cases manually	Natural language query → cited, reasoned answer
Citation trust	No citations; makes up case names	Real citations but no AI reasoning	Every claim verified against source
"Is this still good law?"	Cannot answer	Manual check (partial)	Auto-flag: overruled, distinguished, followed
Case summarization	Generic summary; misses ratio	No summarization feature	Structured: facts, issues, holdings, ratio decidendi
Price	$20/mo (no Indian law depth)	INR 30K–80K/yr	Free tier + INR 999/mo

The Magic Moment

A lawyer types: "Find recent SC cases where bail was denied under PMLA for amounts exceeding 100 crore."

In 8 seconds, they get: 5 relevant cases, each with a 3-line summary, the exact paragraph where bail was discussed, a citation status indicator (green = good law, red = overruled), and a "Copy Citation" button in proper legal format.

The moment they verify that the first citation is real and says exactly what we claim — they're hooked.

User Journey

1. **Sign up** (Google auth, 30 seconds). No credit card for free tier.
2. **Ask a question** in plain English or legal terminology.
3. **Get a reasoned answer** with [1][2] inline citations.
4. **Click any citation** → see the full judgment, highlighted at the relevant paragraph.
5. **Check citation status** → green (good law), yellow (distinguished), red (overruled).
6. **Save research** to a workspace. Export as PDF for filing.
7. **Hit daily limit** → prompted to upgrade to Solo (INR 999/mo).

Real-World Usage Scenarios

1. **Litigation associate preparing for hearing:** Needs to find 3–5 cases supporting anticipatory bail for their client. Currently spends 4+ hours on SCC Online. LawAI: 5 minutes.
2. **Solo practitioner in Lucknow:** Handling a cheque bounce case (Section 138 NI Act). Needs to know the latest quantum of compensation courts are awarding. Can't afford SCC Online. LawAI free tier: instant.
3. **In-house counsel at a startup:** Board wants to know if their new data sharing agreement complies with DPDP Act 2023. Needs statute + recent interpretations. LawAI: structured answer with exact sections cited.

Phase 3: Data Strategy

Data Sources and Accessibility

Source	Content	Volume	Access	Quality
sci.gov.in	SC judgments	~80K	Public, free	Clean (English)
indiankanoon.org	Aggregated case law	Millions	Public, scraping	Mixed quality
indiacode.nic.in	Central Acts	~891	Public, free	Structured
HC websites (25)	HC judgments	500K+	Public, varied	Inconsistent format
Gazette of India	Notifications, amendments	Ongoing	Public	PDF-heavy
SCC Online / Manupatra	Enriched case law	Millions	Paid / proprietary	Gold standard
State legislatures	State-specific acts	Thousands	Hard to access	Fragmented
Tribunal orders	NCLT, ITAT, SAT, etc.	100K+	Semi-public	Messy PDFs

Data Risk: The Quality Gap

Public data (Indian Kanoon, court websites) is free but messy: inconsistent formatting, OCR errors in scanned PDFs, missing metadata, no structured citation links. The gold-standard enriched data (SCC Online, Manupatra) is proprietary. **We must build our own enrichment pipeline** — this is expensive in engineering time but creates a defensible moat.

Data Pipeline Design

1. Ingestion — Scrapers for SC, 25 HCs, Indian Kanoon, India Code, Gazette
Daily cron jobs for new content

↑

2. Extraction — PyMuPDF (PDF), Trafilatura (HTML)
OCR fallback for scanned judgments
(Tesseract / Azure Document Intelligence)

↑

3. Structural Parsing — Extract: parties, judge, date, court, bench size, headnotes, ratio decidendi, citations referenced, statutes interpreted, holdings

↑

4. Chunking + Embedding — Semantic chunks (500–1000 tokens, overlapping)
Voyage AI → 1024-dim vectors → stored in Qdrant Cloud

↑

5. Citation Graph — Forward/backward citations, treatment analysis (followed/distinguished/overruled), statute–case mapping

RAG Architecture (High-Level)

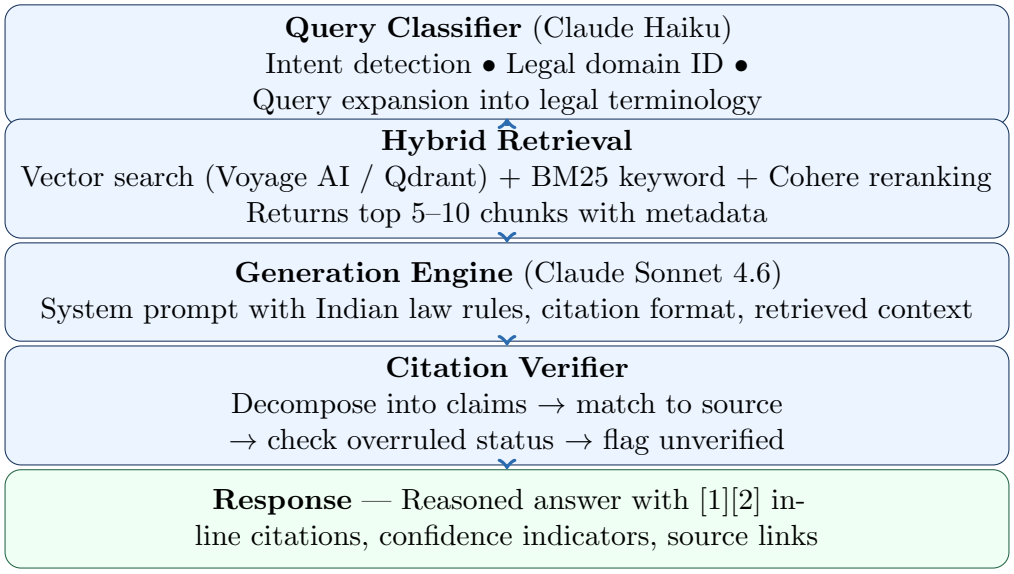

Data Risks

Risk	Description	Mitigation
Accuracy	Messy PDFs, OCR errors, wrong metadata extraction	Multi-pass verification; sample-audit 5% of processed docs with lawyer
Bias	Over-representation of SC/Delhi HC; under-representation of smaller HCs	Explicit coverage tracking; prioritize HCs by case volume
Incompleteness	State laws, tribunal orders, district court judgments missing	Start with SC + top 10 HCs; expand coverage iteratively; be transparent about coverage
Staleness	New judgments, amendments, ordinances	Daily automated scraping + ingestion pipeline
Copyright	Headnotes are copyrighted (Eastern Book Co); raw judgments are not	Only use raw judgment text; build our own headnotes/summaries

Phase 4: AI/ML System Design

Base Model Strategy

Approach	Pros	Cons	Our Decision
API (Claude / GPT)	Best reasoning quality, no infra to manage, fast iteration	Per-token cost, vendor dependency, data leaves your infra	**YES — MVP**
Open-source (Llama, Mistral)	Full control, lower marginal cost at scale, on-prem	Weaker reasoning, expensive to host, slower iteration	Later (Phase 4+)
Fine-tuning	Domain-specific performance (Harvey's 97% preference)	$5M+ for Harvey-scale; hard to update; loses citation ability	Later (1K+ users)
Prompt + RAG	Cheapest, fastest to build, citations possible	Depends on retrieval quality; limited by context window	**YES — core approach**

Our Architecture Decision

Prompt + RAG first, API-based (Claude Sonnet 4.6). This gives us the best reasoning quality at the lowest upfront cost. Fine-tuning is Phase 4+ only after we have real user data to train on. Harvey tried RAG alone first, found it insufficient, and added fine-tuning — but they had $1B+ to spend. We start lean.

Technology Stack

Layer	Technology	Why This Choice
Frontend	Next.js 14 + Tailwind + shadcn/ui	Consistent with venture stack; fast, SSR
Backend	Python FastAPI	Best ecosystem for ML/NLP; async
LLM (Generation)	Claude Sonnet 4.6	Best legal reasoning; 200K context
LLM (Utility)	Claude Haiku 4.5	Cheap classification + routing
Embeddings	Voyage AI (1024-dim)	Top-tier legal embeddings; affordable
Reranking	Cohere Rerank	Proven quality improvement over raw retrieval
Vector DB	Qdrant Cloud	Purpose-built for vector search; scalable
Relational DB + Auth	Supabase (PostgreSQL)	Consistent with venture stack; RLS, auth
Payments	Razorpay	Indian market standard; INR billing
Hosting	Vercel (frontend) + Railway (backend)	Low-ops, auto-scaling

Key AI Capabilities

1. **Legal Reasoning (not just summarization):** Multi-step chain-of-thought prompting forces the model to reason through legal principles, apply them to facts, and distinguish relevant from irrelevant precedent.

2. **Citation Grounding:** Every factual claim in the output maps to a specific paragraph in a specific judgment. Ungrounded claims are flagged or removed.

3. **Multi-document Reasoning:** Retrieval returns chunks from 5–10 different cases. The generation engine synthesizes across them, noting agreements and conflicts.

4. **Structured Output:** Responses follow a consistent format: Summary → Analysis → Relevant Cases → Statute References → Caveats.

Evaluation: How We Measure Correctness

IndiaLaw Bench (Proprietary Benchmark)

Modeled after Harvey's BigLaw Bench. 500+ legal questions with gold-standard answers written by our lawyer partner.

Dimension	How We Measure	Target
Citation Accuracy	Are cited cases real? Do they say what we claim?	≥95%
Legal Correctness	Is the reasoning sound? (lawyer-graded, blind eval)	≥90%
Completeness	Did we miss key precedents? (recall metric)	≥85%
Hallucination Rate	Claims with no source backing (automated check)	<2%
Recency	Citing current law, not overruled judgments	≥95%

Hallucination Reduction Strategy

Harvey still hallucinates ∼17% of the time. Our multi-layer defense:

1. **Retrieval-grounded generation:** model only sees retrieved documents, not its training data
2. **Claim decomposition:** every sentence checked against source
3. **Citation verification:** regex + database lookup confirms cited cases exist
4. **Confidence scoring:** low-confidence claims shown with warning indicator
5. **Human-in-the-loop:** monthly lawyer review of random sample

Biggest Technical Risks

Risk	Impact	Mitigation
Retrieval quality	Wrong cases returned = wrong legal advice	Extensive retrieval benchmarking; hybrid search; reranking
Citation hallucination	Lawyer cites fake case in court = sanctions	Multi-layer verification; "verified" badge only on confirmed citations
Claude API cost spikes	Unit economics collapse at scale	Aggressive caching; Haiku for simple queries; monitor cost per query
Vendor lock-in (Anthropic)	Price increase or API changes	Abstract LLM layer; test with open-source models periodically

Phase 5: Product UX

Interface Decision: Chat + Document Viewer

Not a pure chatbot. Not a pure search engine. A **research workspace** with conversational AI at its center.

Key Screens

1. **Research Chat** — Full-width chat interface. User types a legal question. Response appears with inline [1][2] citations. Each citation is a clickable pill showing case name, court, year.

2. **Citation Panel** (right sidebar) — Click any citation → full judgment opens in a side panel, auto-scrolled to the relevant paragraph. Citation status badge: Good Law / Distinguished / Overruled.

3. **Case Summary View** — Upload a judgment PDF → see structured breakdown: Facts, Issues, Holdings, Ratio Decidendi, Citations Made, Citation Status.

4. **Saved Research** — Organize research by case/matter. Export as formatted PDF with proper legal citations. Share within team (Team plan).

5. **Citation Graph Explorer** — Visual map showing how a case connects to other cases. Which cases cite it? Which cases it cites? What's the treatment chain?

Trust-Building Features (Critical for Lawyers)

- **"Verified Citation" badge:** Only shown when we've confirmed the case exists in our database and the quoted text matches the source.
- **Source transparency:** Every response shows "Based on X documents from [courts]" with links.
- **Confidence indicator:** High/Medium/Low confidence per claim. Low-confidence claims shown in muted text with a caution icon.
- **Legal disclaimer:** Persistent footer: "LawAI is a legal research tool, not a substitute for professional legal advice."
- **"Report Inaccuracy" button:** On every response. Feeds directly into quality improvement pipeline.
- **Coverage transparency:** Show exactly which courts/years are in our database. Don't pretend to cover what we don't.

Phase 6: Go-to-Market Strategy

Who We Sell to First

Not solo practitioners. Not big law. Mid-size litigation firms.

- They have junior associates who waste the most time on research (the pain is acute)
- They bill enough to justify INR 4,999/mo for a team plan
- The senior partner can see immediate ROI (hours saved × billing rate)
- They're big enough to have 5+ lawyers but small enough to not need enterprise sales
- They're in metro cities (Delhi, Mumbai, Bangalore) where tech adoption is highest

Getting the First 10 Customers

Week	Action	Details
1–2	Lawyer partner's network	Personal introductions to 20–30 litigation firms. Demo the prototype in person.
3–4	Bar association events	Present at Delhi/Mumbai bar association meetings. Offer 30-day free trial.
5–6	Legal tech Twitter/LinkedIn	Lawyer partner posts daily showing real research queries and results.
7–8	Law school partnerships	Free access for NLU students. They become advocates for the tool when they join firms.
Ongoing	Referral program	Existing users get 1 free month for every paying referral.

India-Specific Reality: Relationship-Driven Sales

Indian law is a relationship business. Cold emails don't work. LinkedIn ads don't work. What works: a respected lawyer personally vouching for the product to their peers. This is why the lawyer co-founder/advisor is not optional — they *are* the distribution channel for the first 100 customers.

Pricing Model

Tier	Features	Price
Free	10 queries/day, basic research, limited citation depth	INR 0
Solo	Unlimited queries, full citation graph, case upload + summary, export PDF, saved research	INR 999/mo
Team (5 seats)	Everything in Solo + shared workspaces, admin dashboard, team analytics, priority support	INR 4,999/mo
Enterprise	Custom deployment, API access, dedicated support, SLA, SSO	Custom

Why not per-document or per-query? Indian lawyers need predictability. Metered

pricing creates anxiety ("will this query cost me?") and suppresses usage. Flat subscription = unlimited exploration = faster habit formation.

Phase 7: Risks & Failure Modes

Brutal Honesty: Top 10 Ways This Fails

We must confront these head-on. Ignoring them is how startups die.

#	Failure Mode	Probability & Impact	Mitigation
1	**Hallucination causes real harm.** Lawyer cites fake case in court, gets sanctioned, blames LawAI.	High probability, catastrophic impact	Multi-layer verification; "Verified" badge; legal disclaimer; insurance
2	**Lawyers don't trust AI.** SC flagged fake AI citations (Feb 2026). 44% resist new tech.	High probability, high impact	Lawyer co-founder as trust signal; transparency on sources; conservative confidence scoring
3	**Can't compete on data.** Manupatra/SCC have 25 years of enriched data.	Medium probability, high impact	Build citation graph (they don't have this); focus on AI reasoning layer, not raw data volume
4	**Unit economics don't work.** Claude API too expensive at scale; INR 999/mo too cheap.	Medium probability, high impact	Aggressive caching; Haiku routing; raise prices if needed; target team plans over solo
5	**Jhana.ai wins.** Harvard team, $1.8M funded, 10K users, 3-year head start.	Medium probability, medium impact	Differentiate on citation graph + pricing; market is big enough for 2–3 players
6	**Manupatra/SCC build better AI.** They have the data advantage.	Medium probability, medium impact	They're slow (25-year-old orgs); our AI-native architecture is fundamentally different
7	**Regulatory crackdown on AI in law.** BCI issues guidelines restricting AI legal tools.	Low probability, high impact	Position as "research tool for lawyers" not "AI lawyer"; engage with BCI proactively
8	**Free tier users never convert.** 98% stay free, 2% pay. Not enough revenue.	Medium probability, medium impact	Tight free tier (10 queries/day); premium features (citation graph, export) behind paywall
9	**Data pipeline breaks.** Court websites change; scrapers fail; data goes stale.	High probability, low impact	Monitoring + alerts; multiple scraper fallbacks; weekly data freshness audits
10	**Founder distraction.** This is venture #16 in a stack of 18. Insufficient focus.	Medium probability, high impact	Lawyer partner carries domain/GTM; clear milestones; dedicated execution blocks

Phase 8: Execution Plan

Week 1–2: Proof of Concept

- Scrape 10,000 SC judgments (most recent 5 years)
- Build basic document processing pipeline (text extraction + metadata)
- Embed and index in Qdrant
- Build minimal FastAPI endpoint: query → retrieve → generate with Claude
- **Test:** Can we answer "Find SC cases on right to privacy" with real citations?

Month 1: Working Prototype

- Expand to 50K+ SC judgments + top 5 HC judgments
- Add hybrid retrieval (vector + BM25 + reranking)
- Build citation verification layer (check case names against database)
- Basic Next.js frontend (chat + citation links)
- **Test:** Demo to 5 lawyers. Can they use it for real research? What breaks?

Month 3: Beta-Ready Product

- Full coverage: SC + top 10 HCs + Central Acts
- Citation graph v1 (forward/backward citations + basic treatment analysis)
- Case upload + structured summarization
- Auth (Supabase) + usage tracking + free tier limits
- IndiaLaw Bench v1 (200 questions, lawyer-graded)
- **Test:** Beta launch with 50 lawyers. Measure: daily active usage, query success rate, NPS.

Month 6: Paid Launch

- Razorpay payments integration
- IndiaLaw Bench expanded to 500+ questions
- Citation graph v2 (treatment analysis: followed/distinguished/overruled)
- Daily update pipeline operational
- Team workspaces for firm plans
- **Target:** 100–200 paying users. INR 1–2L MRR.

What NOT to Build Early

- × Contract drafting / document generation (complex, liability-heavy)
- × Hindi / regional language support (adds 3–6 months of work)
- × Mobile app (web-first; responsive design is sufficient)
- × WhatsApp bot (distraction; do this only after product-market fit)
- × Fine-tuning (premature without user data)
- × Predictive case outcome / judge analytics (speculative, trust-destroying if wrong)
- × e-Filing integration (heavy regulatory / court system dependency)

Phase 9: Final Verdict

Should We Build This?

Yes. Conditionally.

The market is real (1.7M lawyers, 55M pending cases, $29.5M legal AI market growing 23%). The timing is right (781% funding surge, incumbents are slow, AI capabilities have crossed the quality threshold). The unit economics work (break-even at 300–400 users).

But only if:

1. The lawyer partner is committed (advisor minimum, co-founder ideal)
2. We ship a working prototype in 4 weeks, not 4 months
3. We obsess over citation accuracy above all else — one fake citation and trust is dead
4. We resist the urge to build everything (no drafting, no Hindi, no mobile app in MVP)

What Version of This Idea Actually Wins?

Not "AI legal assistant." Not "ChatGPT for Indian law." Not "legal research tool."

The Winning Version

India's citation intelligence platform.

The product that tells every Indian lawyer, for every case they're about to cite: *Is this still good law? What supports it? What contradicts it? What should you cite instead?*

This is the Shepard's Citations that India never built. It's not just a feature — it's the entire wedge. Everything else (chat, summarization, research) is the AI layer on top of this citation graph.

The ONE Wedge That Dominates

The Indian Citation Graph

Forward citations • Backward citations • Treatment analysis • Statute–case mapping

Nobody has built this. Manupatra has fragments. SCC has fragments. We build the complete, AI-native version.

Every other feature flows from this. Legal research quality depends on it. Trust depends on it. Competitive moat depends on it. If we build nothing else but the citation graph + a simple query interface on top, we have a product worth paying for.

The Decisive Metric

If a lawyer can type a question, get an answer, click a citation, see it's real, and confirm it says what we claim — in under 30 seconds — we win.

Everything we build serves this moment. Everything that doesn't serve this moment is a distraction.

Appendix A: Competitive Intelligence

Harvey AI ($11B Valuation)

- RAG + fine-tuning hybrid; custom 10B-token case law model with OpenAI
- Cascading LLM pipeline; BigLaw Bench for evaluation
- $1B+ raised; $190M ARR; still hallucinates ∼17% of queries
- Pricing: $1,200/seat/month, 20-seat minimums ($288K/yr minimum)
- GTM: top-down enterprise (A&O, PwC, 28% of Am Law 100)

Litera (Lito AI Agent)

- 30-year legal tech company; rules + LLM hybrid architecture
- Free AI tier drove 10× user growth in 3 months
- Multi-doc analysis grid; 11 chainable legal skills
- Native MS 365 integration; 15,000+ firms already using Litera products

Indian Competitors

Player	Funding	Key Strength	Key Weakness
Jhana.ai	INR 15Cr	16M judgments, 10K users, 150+ judges	No citation graph; INR 2.58Cr revenue after 3 years
CaseMine	Undisclosed	CaseIQ + AMICUS; India/US/UK	Shallow; no treatment analysis; generic AI
BharatLaw AI	Undisclosed	Audio digests, voice, multilingual	Limited reasoning depth; no verification
LegitQuest	Undisclosed	400M+ records; enterprise tools	Quality unclear; complex UX; enterprise-only
Manupatra	Established	25 years of enriched data; AI Search (2025)	Bolt-on AI; dated UX; opaque pricing
SCC Online	Established	Gold-standard citations; AI Pro (Azure)	INR 30–80K/yr; enterprise-only

Appendix B: Cost Budget

MVP Build (Months 1–6)

Category	Item	Monthly	6-Mo Total
AI/ML	Claude API (Sonnet 4.6)	$300–500	$2,400
	Voyage AI embeddings	one-time	$100
	Cohere Rerank API	$50	$300
Infrastructure	Qdrant Cloud	$75	$450
	Supabase (PostgreSQL + auth)	$25	$150
	Vercel Pro	$20	$120
	Railway (backend)	$20	$120
Data	Scraping compute	$50–100	$450
Tools	Domain, monitoring, misc	$30	$180
TOTAL		**$600–800/mo**	**$4,270**

Post-Launch Unit Economics

Users	Claude API	Infra	Updates	Total/mo	Revenue/mo
100	$400	$200	$20	**$620**	~$1,200
500	$1,500	$400	$30	**$1,930**	~$6,000
1,000	$3,000	$600	$50	**$3,650**	~$12,000
5,000	$12,000	$1,500	$100	**$13,600**	~$60,000

Break-even: 300–400 paying users. Gross margins improve at scale due to caching and model routing.

Sources

- **Harvey AI:** OpenAI, Contrary Research, ZenML, TechCrunch, CNBC
- **Litera:** Lito Product, LawNext, 10x Growth
- **Hallucinations:** Harvey Study, BigLaw Bench
- **Indian Market:** Grand View Research, Tracxn, Inc42
- **Digitization:** PIB / e-Courts, Bar & Bench
- **Pending Cases:** CS Monitor, NJDG
- **Regulation:** Artificial Lawyer, India AI Summit
- **Competitors:** Jhana.ai, CaseMine, BharatLaw AI, LegitQuest, Manupatra, SCC Online